

Operating Systems

Stack and Context Switch

Giorgio Grisetti

`grisetti@diag.uniroma1.it`

Department of Computer Control and Management Engineering
Sapienza University of Rome

Context of a Process

Process: running program

Context:

- CPU registers
- Process Memory
 - Stack (.stack)
 - Program code (.text), typically read only during execution
 - Initialized Variables (.data)
 - Global Uninitialized Variables (.bss)

If each process uses only its own memory, the execution can be stopped and later recovered by saving/restoring the CPU registers

Coroutines

Coroutine: piece of program that can be “jumped in” and “out”

- In assembly “jmp”.
- In C: need to preserve consistency of the stack.
- Ucontext: portable C library for user level control of contexts

Ucontext: concepts

`struct ucontext_t`; datatype to store a context

- `ucontext_t *uc_link`: pointer to the context that will be resumed when this context returns
- `sigset_t uc_sigmask`: the set of signals that are blocked when this context is active
- `stack_t uc_stack`: the stack used by this context
- `mcontext_t uc_mcontext`: a machine-specific representation of the saved context

getContext

```
int getContext(ucontext_t *ucp) ;
```

- Saves the current context in ucp.
- A subsequent call to setcontext(ucp) will result in the flow of the program continuing from the instruction following setcontext(ucp);

```
struct ucontext_t ctx; ← Here we store our jump point
```

```
int f2() {  
    setContext(&ctx);  
}
```

```
int f1() {  
    ...  
    getContext(&ctx); ← set ctx to jump to the next instruction after  
    ...               getContext  
    f2();  
}
```

setContext

`int setcontext(const ucontext_t *ucp)`

- Sets the current context to ucp, a context that was previously saved.
- The flow will continue from the instruction following the
- `getcontext(ucp)` call issued when SAVING the context

```
    struct ucontext_t ctx;

    int f2() {
        setContext(&ctx);    ← We jump to a saved context [1]
    }

    [1] int f1() {
        ...
        getContext(&ctx);
        ▼ f2();
    }
```

makeContext

```
void makecontext(ucontext_t *ucp, void  
(*func)(), int argc, ...);
```

- creates a trampoline context for function func.
- the context is initialized so that when jumping to it it will start executing the function func
- ucp should have the stack and the signal mask already set before calling makecontext

```
struct ucontext_t ctx;
```

```
int f2(){  
    ...  
}
```

```
int main(){
```

```
    ....
```

```
    makeContext(&ctx, f2,...);
```

← ctx initialized so that when called starts f2

```
    ...
```

```
    setContext(&ctx);
```

← flow continues from f2

```
}
```

swapContext

```
int swapcontext(ucontext_t *oucp,  
const ucontext_t *ucp) ;
```

- saves the current context in oucp, and jumps to ucp

Full example

```
ucontext_t main_context, f1_context, f2_context;

void f1(){
    printf("f1 started\n");
    for (int i=0; i<num_iterations; i++) {
        printf("f1: %d\n", i);
        swapcontext(&f1_context, &f2_context);
    }
    setcontext(&main_context);
}

void f2(){
    printf("f2 started\n");
    for (int i=0; i<num_iterations; i++) {
        printf("f2: %d\n", i);
        swapcontext(&f2_context, &f1_context);
    }
    setcontext(&main_context);
}

char f1_stack[STACK_SIZE];
char f2_stack[STACK_SIZE];

int main(){
    //get a context from main
    getcontext(&f1_context);

    // set the stack of f1 to the right place
    f1_context.uc_stack.ss_sp=f1_stack;
    f1_context.uc_stack.ss_size = STACK_SIZE;
    f1_context.uc_stack.ss_flags = 0;
    f1_context.uc_link=&main_context;

    // create a trampoline for the first function
    makecontext(&f1_context, f1, 0, 0);

    // always remember to initialize
    // a new context from something known
    f2_context=f1_context;

    f2_context.uc_stack.ss_sp=f2_stack;
    f2_context.uc_stack.ss_size = STACK_SIZE;
    f2_context.uc_stack.ss_flags = 0;
    f2_context.uc_link=&main_context;

    // create a trampoline for the second function
    makecontext(&f2_context, f2, 0, 0);

    // this passes control to f2.
    // and saves the current context in main_context

    swapcontext(&main_context, &f1_context);
    // we will jump back here

    printf("exiting\n");
}
```

Exercise

- Modify the program above to spin on 10 different contexts instead of two

Preemptive multitasking on AVR

We want to implement an timer controlled preemptive task switcher on our arduino.

- **Task Control Blocks:** stored in double linked list
- Always at least one process in running

Initialization

- Fill in TCB data
- Prepare all stack frames so that the Program Counter stored on the stack points to a launcher for the thread function, and all registers clean

Start

- Change stack pointer to first tcb
- Pull all registers
- Return from function

Context switch (once all is set), on interrupt:

- Save all registers on stack
- Change stack pointer
- Pull all registers from stack
- Return from interrupt

Task Control Block

```
#pragma once
#include <stdint.h>
#include <stddef.h>

#define OK      0
#define ERROR -1

typedef uint8_t* Pointer;
typedef void (* ThreadFn) (uint32_t thread_arg);

typedef enum {Running=0x0, Terminated=0x1, Ready=0x2} ThreadStatus;

// thread control block
typedef struct TCB {
    Pointer sp_save_ptr;

    ThreadFn thread_fn;
    uint32_t thread_arg;

    struct TCB* next;
    struct TCB* prev;

    Pointer stack_bottom;          /* Pointer to bottom of stack allocation */
    uint32_t stack_size;          /* Size of stack allocation in bytes */
    ThreadStatus status;
} TCB;

void TCB_create(TCB* tcb, Pointer stack_top, ThreadFn thread_fn, uint32_t thread_arg);
```

TCB Create

```
void TCB_create(TCB* tcb, Pointer stack_top, ThreadFn thread_fn, uint32_t thread_arg){
    //initialize variables
    tcb->thread_fn=thread_fn;
    tcb->thread_arg=thread_arg;
    tcb->prev=NULL;
    tcb->next=NULL;
    tcb->status=Ready;

    /** prepare stack for process **/

    uint8_t *stack_ptr = (uint8_t *)stack_top;
    //write the return address of the function being called (the trampoline)
    *stack_ptr-- = (uint8_t)((uint16_t)_trampoline & 0xFF);
    *stack_ptr-- = (uint8_t)(((uint16_t)_trampoline >> 8) & 0xFF);
    *stack_ptr-- = 0; // store an additional segment register (atMega2560)

    /**
     * Store starting register values for R2-R17, R28-R29
     */
    *stack_ptr-- = 0x00;    /* R2 */
    *stack_ptr-- = 0x00;    /* R3 */
    ..... // here we save all other registers.....
    *stack_ptr-- = 0x00;    /* R28 */
    *stack_ptr-- = 0x00;    /* R29 */
    *stack_ptr-- = 0x00;    /* RAMPZ */
    *stack_ptr-- = 0x00;    /* EIND */

    // store stack pointer
    tcb->sp_save_ptr = stack_ptr;
}
```

TCB Create, trampoline

- The trampoline is a convenient function without parameters that calls the function whose pointer is stored in the `current_tcb` global variable
- Not to mess up with calling conventions ;-)

```
static void _trampoline(void) {
    sei();
    /* Call the thread entry point */
    if (current_tcb && current_tcb->thread_fn) {
        (*current_tcb->thread_fn)(current_tcb->thread_arg);
    }

    // set the thread to terminated, when the above function finishes
    current_tcb->status=Terminated;
}
```

TCB Queue

The TCBs are stored in a double linked list
No memory allocation

- Two actions:
 - Take out the element at the beginning of the list
 - Put an element out of the list at its tail

```
// simple double linked list of TCBs
typedef struct {
    struct TCB* first;
    struct TCB* last;
    uint8_t size;
} TCBList;

// global list of tcbs containing the running processes
extern TCBList tcb_queue;

// removes (if any) first tcb from the list
TCB* TCBList_dequeue(TCBList* list);

// adds new detached tcb to the list
uint8_t TCBList_enqueue(TCBList* list, TCB* tcb);
```

Context Switch

```
//void archContextSwitch (ATOM_TCB *old_tcb_ptr, ATOM_TCB *new_tcb_ptr)

.global archContextSwitch
archContextSwitch:

    /**
     * Parameter locations:
     *  old_tcb_ptr = R25-R24
     *  new_tcb_ptr = R23-R22
     */

    /**
     * Save registers R2-R17, R28-R29.
     */
    push r2
    .....
    push r29
    // save RAMPZ and EIND
    in r0,_SFR_IO_ADDR(RAMPZ)
    push r0
    in r0,_SFR_IO_ADDR(EIND)
    push r0

    // Save the final stack pointer to the TCB.
    in r16,_SFR_IO_ADDR(SPL)
    in r17,_SFR_IO_ADDR(SPH)
    mov r28,r24
    mov r29,r25
    st Y,r16
    std Y+1,r17

    //get SP from new TCB
    mov r28,r22
    mov r29,r23
    ld r16,Y
    ldd r17,Y+1

    // switch stack
    out _SFR_IO_ADDR(SPL),r16
    out _SFR_IO_ADDR(SPH),r17

    // restore status
    pop r0
    in r0,_SFR_IO_ADDR(EIND)
    pop r0
    in r0,_SFR_IO_ADDR(RAMPZ)
    pop r29
    ...
    pop r2
    ret
```


First Thread Restore

Is just the bottom part of the context switch

```
void archFirstThreadRestore (ATOM_TCB *new_tcb_ptr)
*/
.global archFirstThreadRestore

archFirstThreadRestore:

    /**
     * Parameter locations:
     * new_tcb_ptr = R25-R24
     */

    //get SP from new TCB
    mov r28,r24
    mov r29,r25
    ld r16,Y
    ldd r17,Y+1

    // switch stack
    out _SFR_IO_ADDR(SPL),r16
    out _SFR_IO_ADDR(SPH),r17

    // restore status
    pop r0
    in r0,_SFR_IO_ADDR(EIND)
    pop r0
    in r0,_SFR_IO_ADDR(RAMPZ)
    pop r29
    ...
    pop r2
    ret
```

Schedule

The final scheduler consists of:

- The current process, and the head of a list of thread control blocks
- Two functions:
 - startSchedule (initializes timers, and gives control to first thread)
 - schedule (called in the timer interrupt), that switches context

```
TCB* current_tcb=NULL;
```

```
// the running queue
TCBList running_queue={
    .first=NULL,
    .last=NULL,
    .size=0
};
```

```
void startSchedule(void) {
    cli();
    current_tcb=TCBList_dequeue(&running_queue);
    assert(current_tcb);
    timerStart();
    archFirstThreadRestore(current_tcb);
}
```

```
void schedule(void) {
    TCB* old_tcb=current_tcb;
    // we put back the current thread in the queue
    TCBList_enqueue(&running_queue, current_tcb);

    // we fetch the next;
    current_tcb=TCBList_dequeue(&running_queue);
    // we jump to it
    //(useless if it is the only process)
    if (old_tcb!=current_tcb)
        archContextSwitch(old_tcb, current_tcb);
}
```

Run, baby run

```
TCB idle_tcb;
uint8_t idle_stack[IDLE_STACK_SIZE];
void idle_fn(uint32_t thread_arg){
    while(1) {
        cli();
        printf("i\n");
        sei();
        _delay_ms(10);
    }
}

TCB p1_tcb;
uint8_t p1_stack[THREAD_STACK_SIZE];
void p1_fn(uint32_t arg ){
    while(1){
        cli();
        printf("p1\n");
        sei();
        _delay_ms(10);
    }
}

TCB p2_tcb;
uint8_t p2_stack[THREAD_STACK_SIZE];
void p2_fn(uint32_t arg ){
    while(1){
        cli();
        printf("p2\n");
        sei();
        _delay_ms(10);
    }
}
```

```
int main(void){
    // we need printf for debugging
    printf_init();

    TCB_create(&idle_tcb,
               idle_stack+IDLE_STACK_SIZE-1,
               idle_fn,
               0);

    TCB_create(&p1_tcb,
               p1_stack+THREAD_STACK_SIZE-1,
               p1_fn,
               0);

    TCB_create(&p2_tcb,
               p2_stack+THREAD_STACK_SIZE-1,
               p2_fn,
               0);

    TCBLIST_enqueue(&running_queue, &p1_tcb);
    TCBLIST_enqueue(&running_queue, &p2_tcb);
    TCBLIST_enqueue(&running_queue, &idle_tcb);

    printf("starting\n");
    startSchedule();
}
```