Operating Systems

CPU Scheduling

Giorgio Grisetti

grisetti@diag.uniroma1.it

Department of Computer Control and Management Engineering
Sapienza University of Rome

CPU Scheduler

Selects which processes to grant the cpu among those in the ready queue

Invoked when

1. a process switches from running to waiting

2. terminates

3. switched from running to ready

4. switches from waiting to ready

preemptive

non-preemptive

CPU Scheduling

In a multiprogrammed system, the CPU scheduler is always invoked upon an I/O request by a process

- during I/O the CPU will be useless for the process requesting it.
- wisdom suggests to assign the CPU to another process (if any) in ready.

Non Preemptive (batch) schedulers

a running process is never evicted from CPU if it does not request an I/O or terminates.

Preemptive

a running process can be evicted and put in the ready queue(s) before it requests an I/O, since it has consumed its CPU time quantum

Dispatcher

Module that gives the CPU to the process selected by the short term scheduler.

Actions:

- switch to kernel mode and save current process state (see preamble)

- switch to user mode while restoring new process state (see postamble)
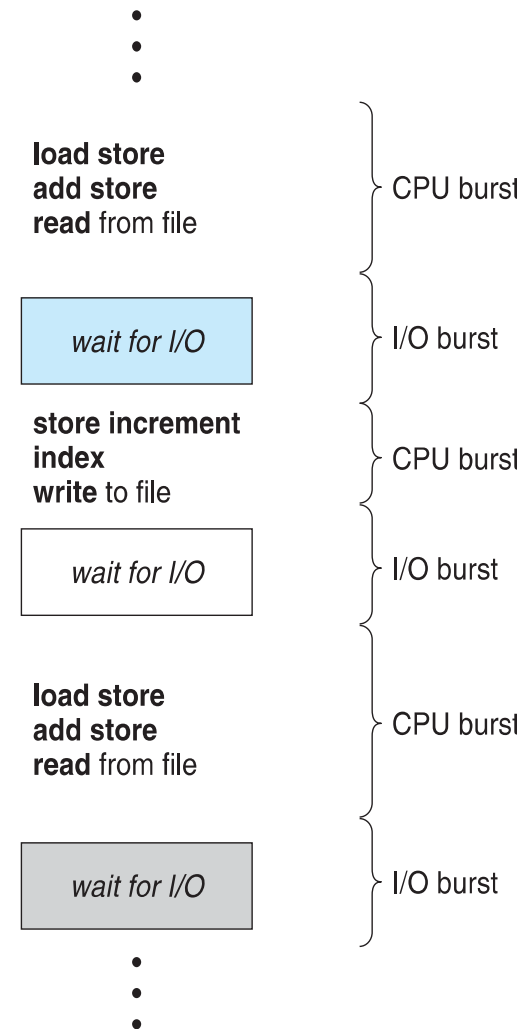
Switching costs time:

- Saving/Restoriing CPU state

- **Cache after a task switch is invalid**

CPU/IO Bursts

User programs are characterized by an alternation of

- CPU bursts
  - time interval where the CPU is used
  - CPU is the bottleneck
  - I/O rests
- I/O bursts
  - time interval where the I/O is used
  - I/O is the bottleneck
  - CPU rests

CPU burst distribution is of main concern

⋮

|	
load store **add store** **read** from file	CPU burst
wait for I/O	I/O burst
store increment **index** **write** to file	CPU burst
wait for I/O	I/O burst
load store **add store** **read** from file	CPU burst
wait for I/O	I/O burst

⋮

CPU Burst Distribution

Is a distribution over the length of cpu bursts.

Characterizes the CPU behavior of a program.

Can be computed by calculating an histogram while running the program

- on the x: duration of cpu burst (interval)

- on the y: # of times an interval of duration x occurs
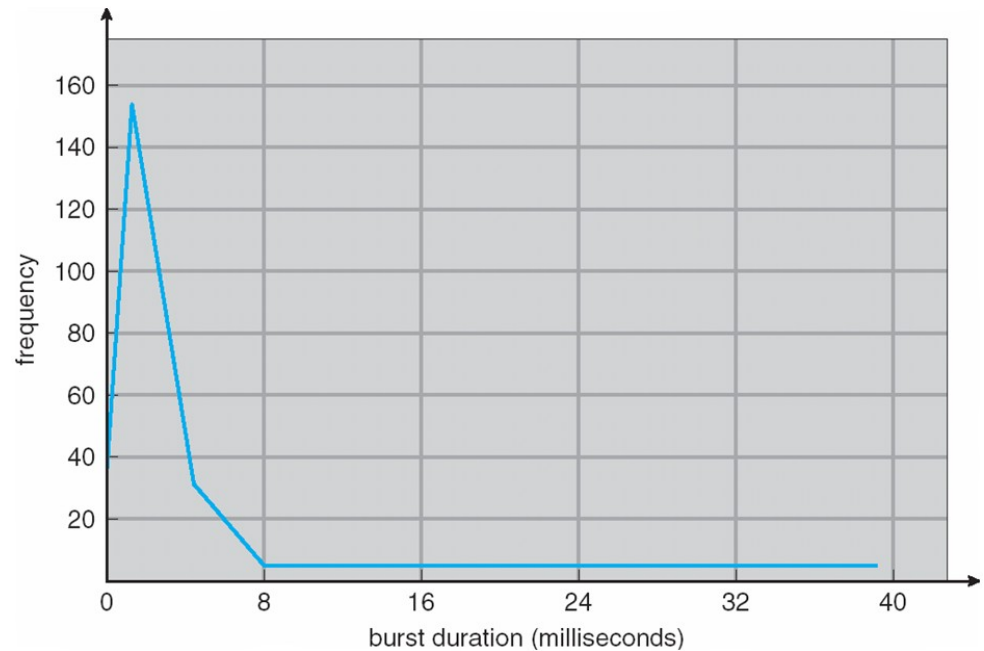
CPU Scheduling: Metrics

The choice of which process to give the CPU next affects the behavior of the entire system.

The behavior can be monitored by the following indicators:

- **CPU utilization**: fraction of the time the CPU is used by the processes.

- **Turnaround time**: time to complete a process.

- **Throughput**: number of processes that complete in a time unit.

- **Waiting time**: how long a process has been waiting in the ready queue.

- **Response time:** how long does it take for a process that receives a command to start providing the answer

CPU Scheduling: Metrics

The choice of which process to give the CPU next affects the behavior of the entire system.

The behavior can be monitored by the following indicators:

MAXIMIZE

- **CPU utilization**: fraction of the time the CPU is used by the processes.

- **Throughput:** number of processes that complete in a time unit.

MINIMIZE

- **Turnaround time**: time to complete a process.

- **Waiting time**: how long a process has been waiting in the ready queue.

- **Response time**: how long does it take for a process that receives a command to start providing the answer

Describing a Process

To the extent of the CPU scheduler, process can be summarized by

- its arrival time
- a list of "actions", each action can be
    - a CPU burst of a given duration
    - an I/O burst on a device, of a given duration

For non interactive processes asynchronous I/O is not considered

Process P1:

- arrival, T=100
- CPU, D=5
- DISK1_IO, D=5000
- CPU, D=1
- DISK2_IO, D=5000
- CPU, D=1
- ....

Process P2:

- arrival, T=10
- CPU, D=50
- DISK1_IO, D=100
- CPU, D=1000
- DISK2_IO, D=100
- CPU, D=1000
- ...

Illustrating a Scheduler

The behavior of a scheduler executing a set of processes can be illustrated through a diagram

- Time on the X axis
- One row per I/O resource*
- One row per CPU core*

The rows of a resource are filled with a color/id corresponding to the process using that resource in that time interval.

The dispatch latencies are usually neglected.

To "paint" these diagrams it is usually helpful to use N rows, one per process.

* for simplicity we will consider 1 core and 1 I/O

First Come First Served (FCFS)

Non Preemptive

Key idea:

- pick the first process that is in ready
- when an I/O terminates, a process moves from waiting to ready, and is but in the back of the queue

Example:

Process	Burst Time	Arrival Time
P_1	24	0
P_2	3	1
P_3	3	2

- The Gantt Chart for the schedule is:

P_1	P_2	P_3
0 24	27	30

- Waiting time for P_1 = 0; P_2 = 23; P_3 = 25
- Average waiting time: $(0 + 23 + 25)/3 = 16$

FCFS

Suppose that the processes arrive in the order:

P_2 , P_3 , P_1

- **GANTT**

P_2	P_3	P_1
0	3	6 30

- Waiting time for P_1 = 6; P_2 = 0; P_3 = 3
- Average waiting time: (6 + 0 + 3)/3 = 3

Much better than previous case

Convoy effect - short process behind long process

Consider one CPU - bound and many I/O-bound processes

Approximating SJF

We can still use the SJF schema, if we have a way to "predict" the next cpu burst of the process.

The same process typically has a "cyclic" behavior

- CPU bursts of the same length, followed by IO burst of similar length
- The behavior might change during the execution of the program, but stays relatively steady for relatively long periods
- e.g. A programming IDE
  - editing: short CPU bursts, long I/O bursts (keyboard is slow)
  - compiling: long burst, medium I/Os (disk is fast)

Shortest Job First (SJF)

Non Preemptive

Idea:

- Select the job in ready whose **next** CPU burst is the shortest

  (keep the ready list sorted by next cpu burst)

Pros:

- Achieves the optimal minimum average waiting time for a given set of processes
- Maximizes throughput

Cons:

- Requires to know the behavior of a process (CPU and I/O bursts) in advance (often not possible in practice)
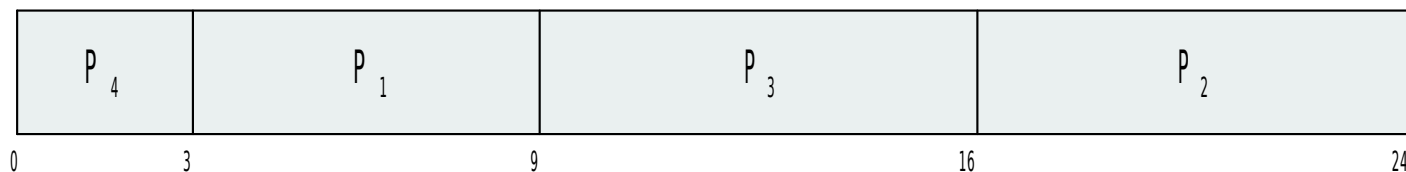
Implications:

To be used at exams when dealing with multiple exercises, provided you have a good estimate of how long each exercise will take

SJF example

Assume all processes arrive at time 0

Process	Burst Time
P_1	6
P_2	8
P_3	7
P_4	3

- SJF scheduling chart

P_4	P_1	P_3	P_2

0   3   9   16   24

- Average waiting time = (3 + 16 + 9 + 0) / 4 = 7
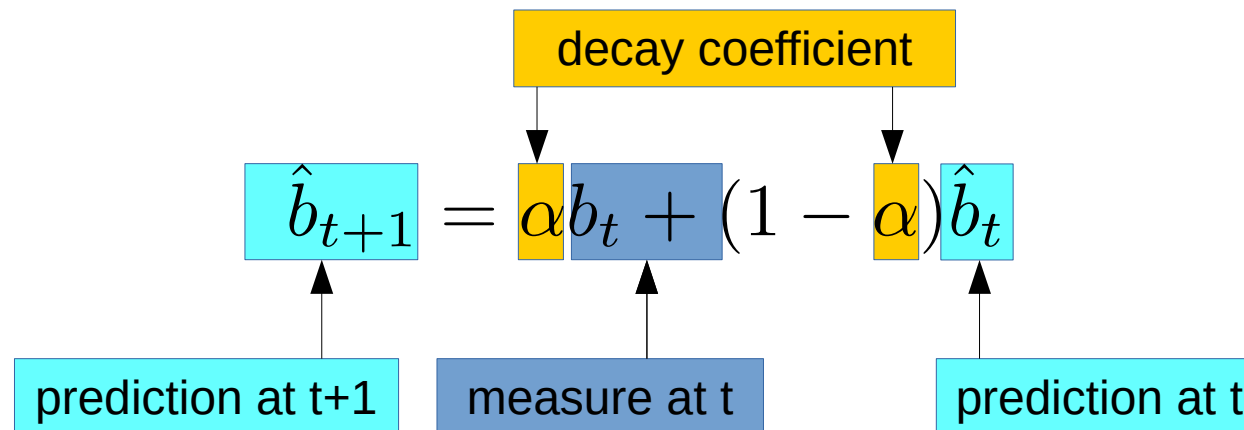
SJF: Predicting Next Bursts

Rule of thumb:

- the next burst will be as long as the current one

Issue:

- sudden "spikes" in the bursts might degrade overall quality of prediction

Solution

- use a discrete low pass filter to smooth the spike (aka exponential mean)

decay coefficient

$$\hat{b}_{t+1} = \alpha b_t + (1 - \alpha)\hat{b}_t$$

prediction at t+1

measure at t

prediction at t

Priority Scheduler

Non-Preemptive

Processes are assigned a priority (int p)

- traditionally if p1<p2, p1 has highest priority

Processes in ready with highest priority are executed first

- SJF is a priority scheduler where the priority is the inverse of the next cpu burst

Issues:

- Starvation: low priority processes might never be executed

Solution:

- Aging: increase the priority as a process spends time in the ready queue

Priority Example

Process	Burst Time	Priority	Arrival Time
P_1	11	3	0
P_2	5	1	1
P_3	2	4	2
P_4	1	5	3

- Priority scheduling Gantt Chart

P_1	P_2	P_1	P_3	P_4
0 1	6	16	18	19

- Average waiting time = 8.2 msec

Preemptive Schedulers

All schemas discussed above can be extended to preemptive schedulers.

The core of a preemptive scheduler is a routine that can put in ready a running process that has not yet requested an I/O.

Each process gets a small unit of CPU time (the cpu quantum q, usually 10/100 ms). If after this time the process is still using the CPU, it is preempted and put in the ready queue.

- **If evicted process is put at the end -> Round Robin Scheme (RR)**

Scheduler is invoked upon
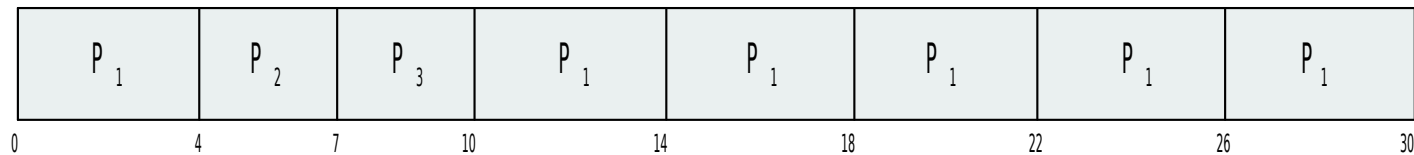
- I/O requests
- timer interrupt

With N processes in ready and a quantum of q, no process waits more than $(N-1)*q$

RR example

Assuming *q*=4

Process	Burst Time
P_1	24
P_2	3
P_3	3

The Gantt chart is:

P_1	P_2	P_3	P_1	P_1	P_1	P_1	P_1

0   4   7   10   14   18   22   26   30

Considerations:

▪Typically, higher average turnaround than SJF, but better **response**

▪*q* should be large compared to context switch time

▪*q* usually 10ms to 100ms, context switch < 10 usec

q and Context Switch Time

- The smaller the quantum, the more the context switches

- Too many context switches might waste CPU

- Usually q chosen so that 80% of cpu bursts are shorter than q

process time = 10

	quantum	context switches
	12	0
	6	1
	1	9
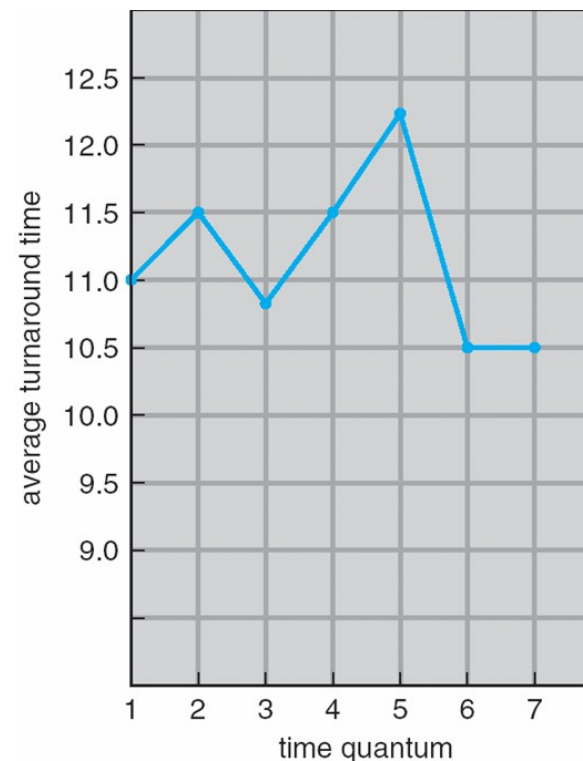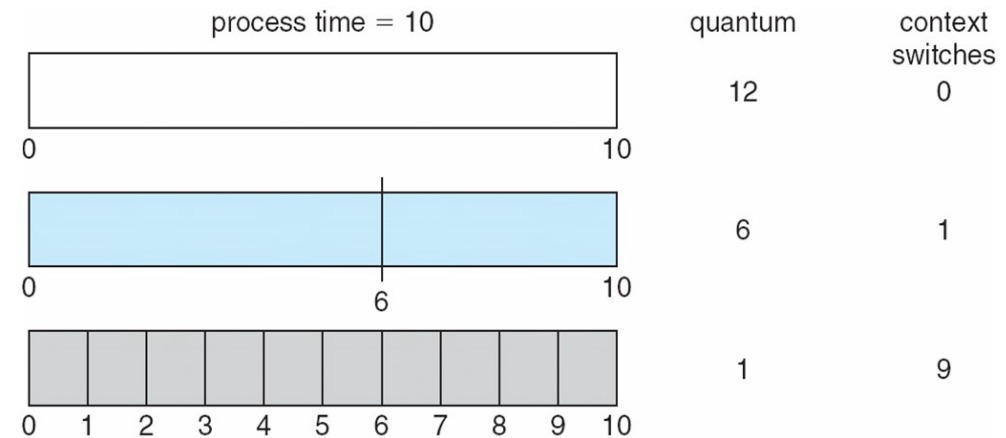
process	time
P_1	6
P_2	3
P_3	1
P_4	7

Matching Application's Needs

The scheduling metric to be optimized depends on the application

- Interactive processes (e.g. editors): response time
- Batch Processes: (e.g. building an application) throughput

The scheduler is responsible of optimizing the parameters for each application.

Idea:

- Threat applications differently

Multilevel Queue

Ready queue is partitioned into separate queues, eg:

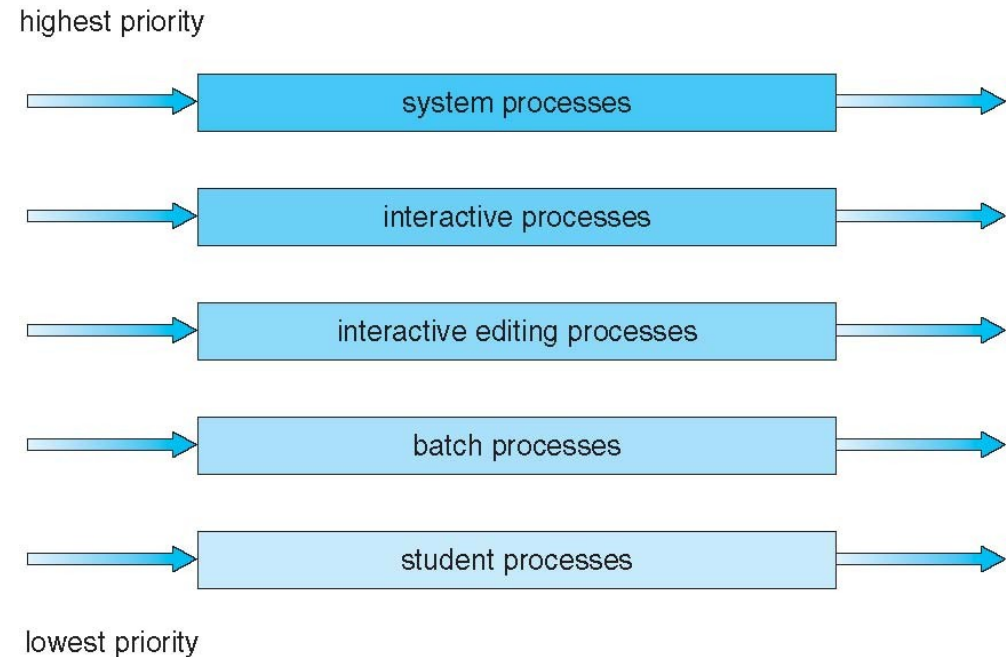- **foreground** (interactive)
- **background** (batch)

Process permanently in a given queue

Each queue has its own scheduling algorithm:
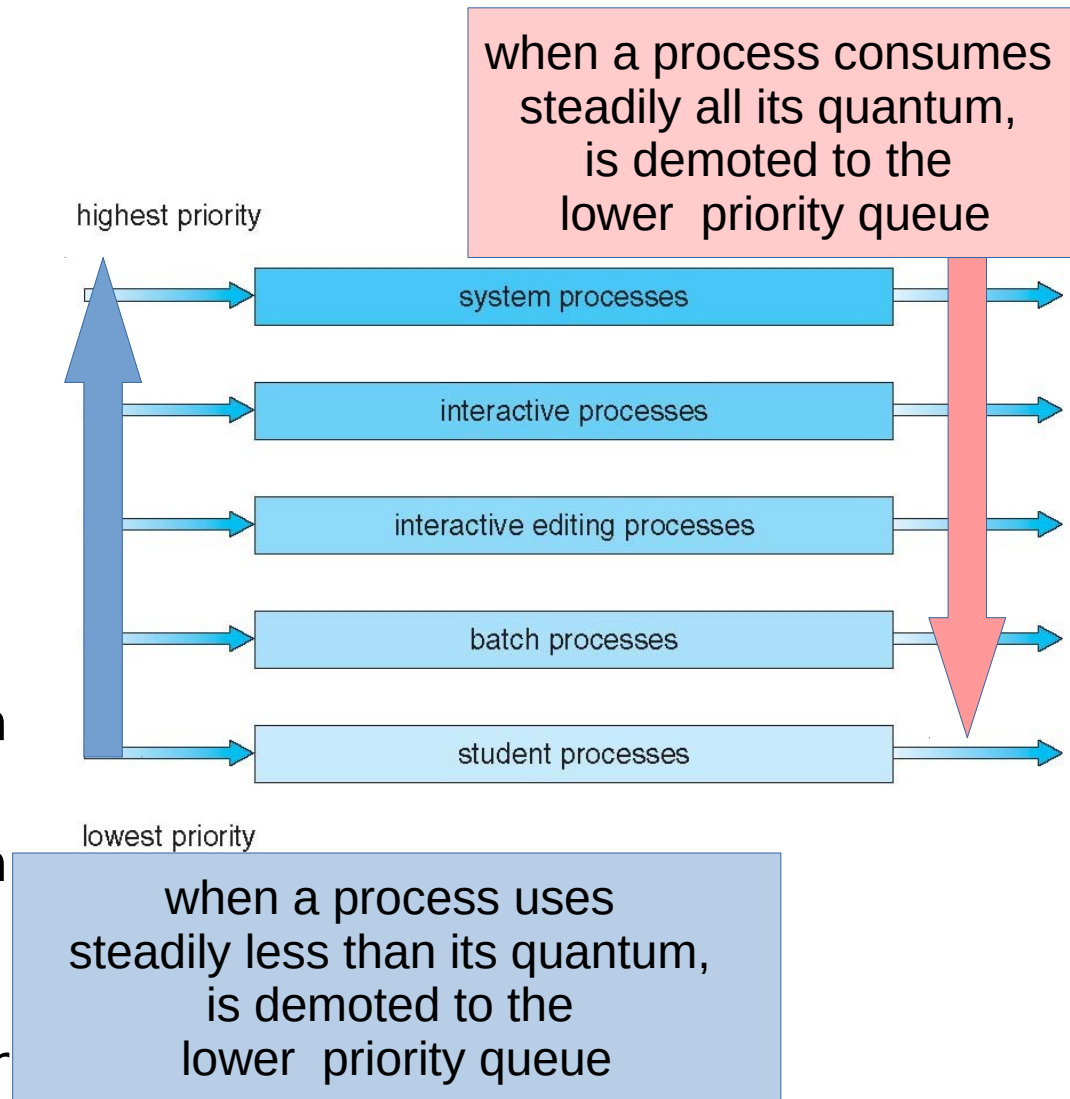
- foreground – RR
- background – FCFS

Scheduling must be done between the queues:

- Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
- Time slice – each queue gets a certain amount of CPU time which is distributed to its processes; i.e., 80% to foreground in RR
- 20% to background in FCFS

highest priority

system processes

interactive processes

interactive editing processes

batch processes

student processes

lowest priority

Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way

- Multilevel-feedback-queue scheduler defined by the following parameters:

  - number of queues

  - scheduling algorithms for each queue

  - method used to determine when to upgrade a process

  - method used to determine when to demote a process

  - method used to determine which queue a process will enter when that process needs service

when a process consumes steadily all its quantum, is demoted to the lower priority queue

highest priority

system processes

interactive processes

interactive editing processes

batch processes

student processes

lowest priority

when a process uses steadily less than its quantum, is demoted to the lower priority queue

Multiple Processor Scheduling

CPU scheduling more complex when multiple CPUs are available

Homogeneous processors within a multiprocessor

Asymmetric multiprocessing – only one processor accesses the system data structures, alleviating the need for data sharing

Symmetric multiprocessing (**SMP**) – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes

- Currently, most common

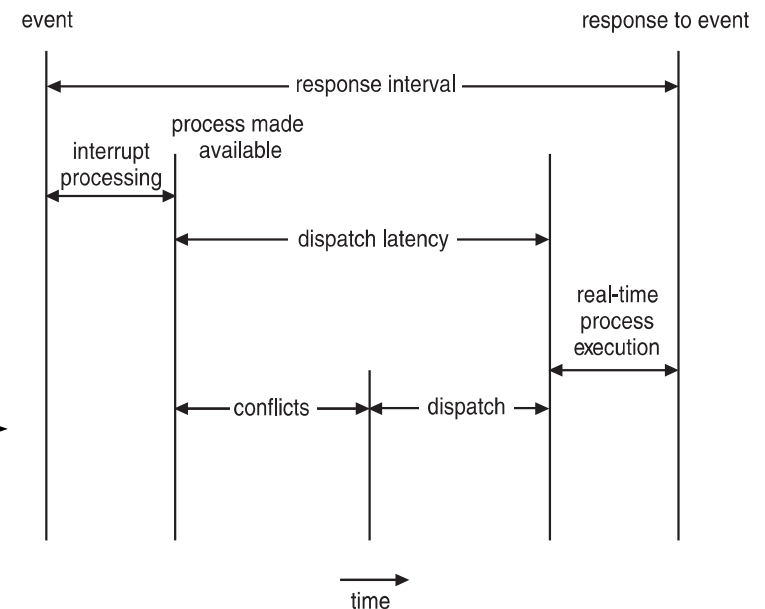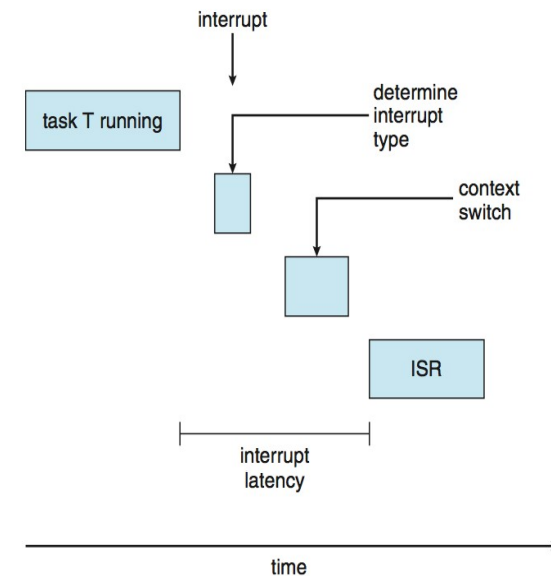**Processor affinity** – process has affinity for processor on which it is currently running

- **soft affinity**
- **hard affinity**
- Variations including **processor sets**

Real Time CPU Scheduling

Soft real-time systems no guarantee as to when critical real-time process will be scheduled

- **Hard real-time systems** task must be serviced by its deadline

- Two types of latencies affect performance

  - Interrupt latency – time from arrival of interrupt to start of routine that services interrupt

  - Dispatch latency – time for schedule to take current process off CPU and switch to another

interrupt

task T running

determine interrupt type

context switch

ISR

interrupt latency

time

event

response to event

response interval

interrupt processing

process made available

dispatch latency

real-time process execution

conflicts

dispatch

time

Real Time Priority Scheduling

For **soft real-time** scheduling, scheduler must support preemptive, priority-based scheduling

For **hard real-time** the scheduler must also provide ability to meet deadlines

- Processes have new characteristics: **periodic** ones require CPU at constant intervals
- Has processing time t, deadline d, period p
- $0 \leq t \leq d \leq p$
- **Rate** of periodic task is $1/p$
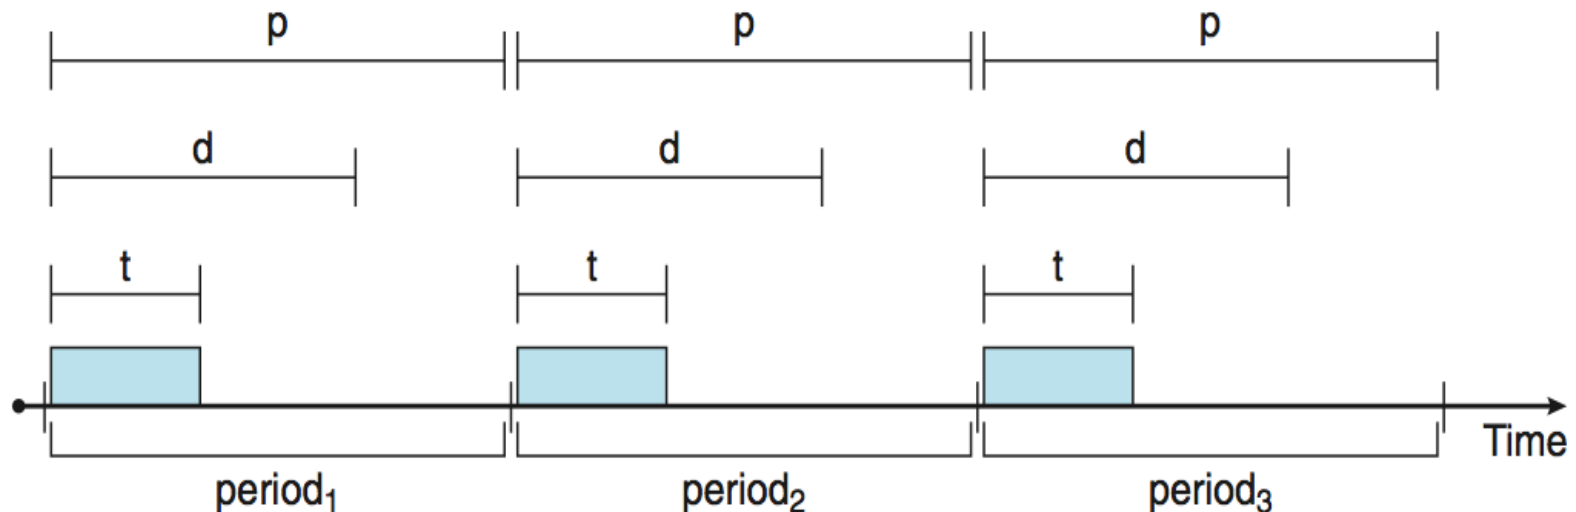
Evaluating a Scheduler

- How to select CPU-scheduling algorithm for an OS?

- Determine criteria, then evaluate algorithms

- **Deterministic modeling**

  - Type of **analytic evaluation**

  - Takes a particular predetermined workload and defines the performance of each algorithm for that workload

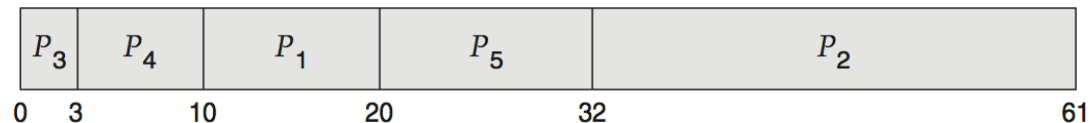- Consider 5 processes arriving at time 0:

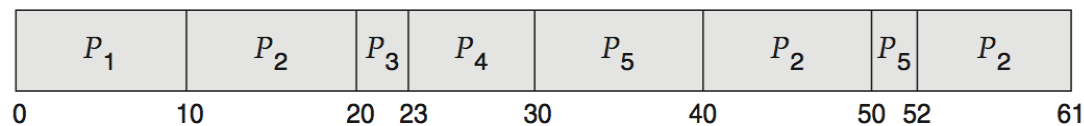| Process	Burst Time
P_1	10
P_2	29
P_3	3
P_4	7
P_5	12

Deterministic Evaluation

- For each algorithm, calculate minimum average waiting time

- Simple and fast, but requires exact numbers for input, applies only to those inputs

  - FCS is 28ms:

P_1	P_2	P_3	P_4	P_5

0   10   39   42   49   61

  - Non-preemptive SFJ is 13ms:

P_3	P_4	P_1	P_5	P_2

0   3   10   20   32   61

  - RR is 23ms:

P_1	P_2	P_3	P_4	P_5	P_2	P_5	P_2

0   10   20   23   30   40   50 52   61

Queuing Models

Describes the arrival of processes, and CPU and I/O bursts probabilistically

- Commonly exponential, and described by mean
- Computes average throughput, utilization, waiting time, etc

Computer system described as network of servers, each with queue of waiting processes

- Knowing arrival rates and service rates
- Computes utilization, average queue length, average wait time, etc

LITTLE's Formula

- Little's law – in steady state, processes leaving queue must equal processes arriving, thus:
  $$n = \lambda \times W$$
  - Valid for any scheduling algorithm and arrival distribution
- n = average queue length
- W = average waiting time in queue
- λ = average arrival rate into queue

- For example, if on average 7 processes arrive per second, and normally 14 processes in queue, then average waiting time per process = 2 seconds

Simulations

Queueing models limited

Simulations more accurate

- Programmed model of computer system
- Clock is a variable
- Gather statistics indicating algorithm performance
- Data to drive simulation gathered via
  - Random number generator according to probabilities
  - Distributions defined mathematically or empirically
  - Trace tapes record sequences of real events in real systems