

Operating Systems

Virtual Memory

Giorgio Grisetti

`grisetti@diag.uniroma1.it`

Department of Computer Control and Management Engineering
Sapienza University of Rome

Facts

Code needs to be in memory to execute, but entire program rarely used

- Error code, unusual routines, large data structures
- Entire program code not needed at same time
- Consider ability to execute partially-loaded program
 - Program no longer constrained by limits of physical memory
 - Each program takes less memory while running -> more programs run at the same time
 - Increased CPU utilization and throughput with no increase in response time or turnaround time
 - Less I/O needed to load or swap programs into memory -> each user program runs faster

Facts

Virtual memory – separation of user logical memory from physical memory

- Only part of the program needs to be in memory for execution
- Logical address space can therefore be much larger than physical address space
- Allows address spaces to be shared by several processes
- Allows for more efficient process creation
- More programs running concurrently
- Less I/O needed to load or swap processes

Virtual address space – logical view of how process is stored in memory

- Usually start at address 0, contiguous addresses until end of space
- Meanwhile, physical memory organized in page frames
- MMU must map logical to physical

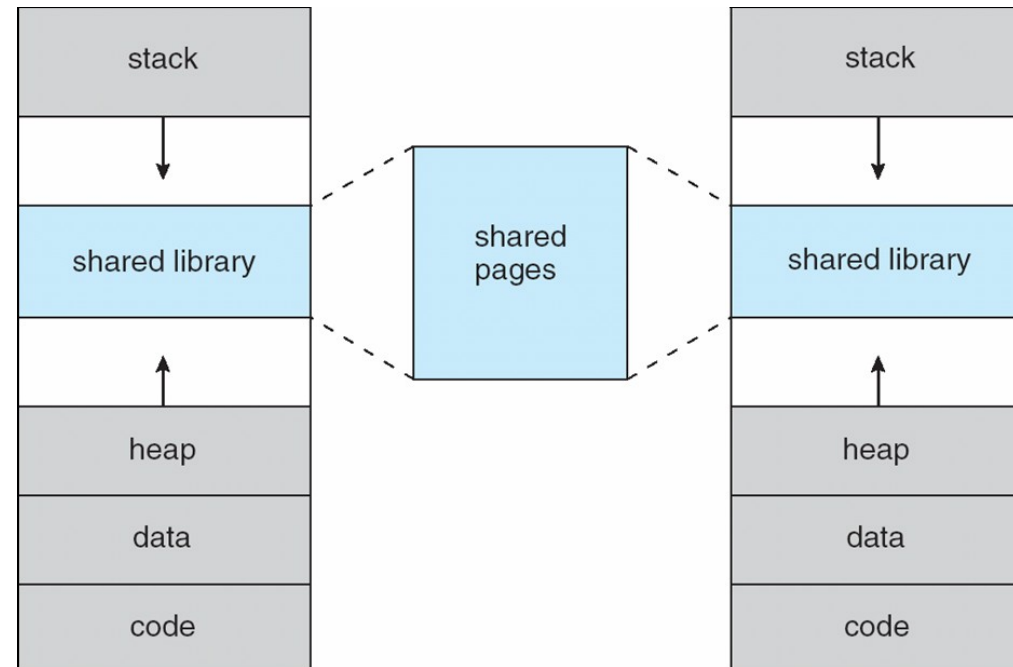
Virtual memory can be implemented via:

- Demand paging
- Demand segmentation

Virtual Address Space

Usually design logical address space for stack to start at Max logical address and grow “down” while heap grows “up”

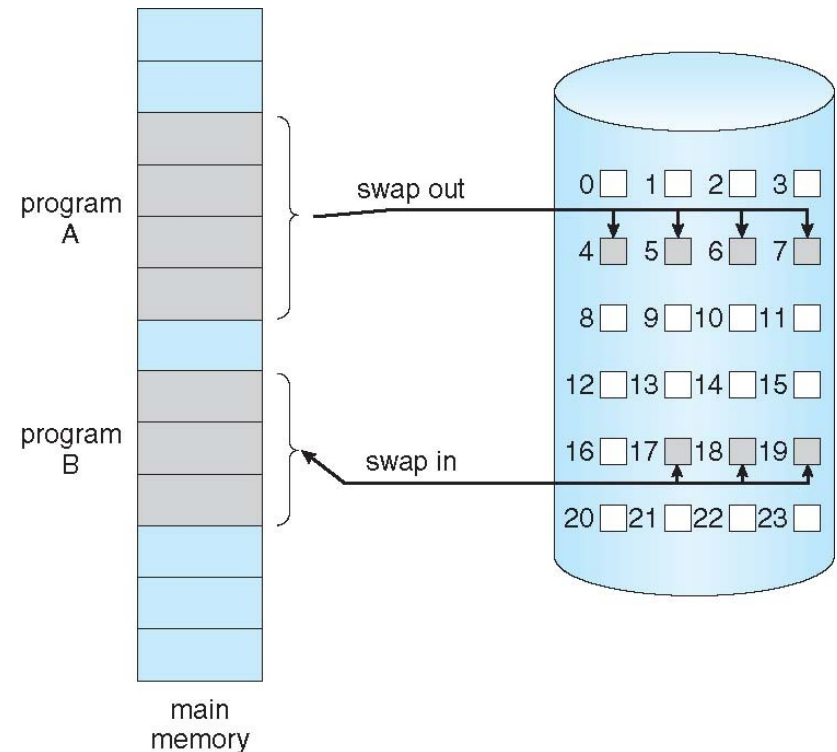
- Maximizes address space use
- Unused address space between the two is hole
- No physical memory needed until heap or stack grows to a given new page
- Enables **sparse** address spaces with holes left for growth, dynamically linked libraries, etc
- System libraries shared via mapping into virtual address space
- Shared memory by mapping pages read-write into virtual address space
- Pages can be shared during `fork()`, speeding process creation



Demand Paging

Instead of bringing the entire process into memory at load time, bring a page into memory only when it is needed

- Less I/O needed, no unnecessary I/O
- Less memory needed
- Faster response
- More users
- Similar to paging system with swapping (diagram on right)
- Page is needed -> reference to it
 - invalid reference -> abort
 - not-in-memory -> bring to memory
- **Lazy swapper** – never swaps a page into memory unless page will be needed
- Swapper that deals with pages is a **pager**



Demand Paging

Pure demand paging: extreme case, start process with *no* pages in memory

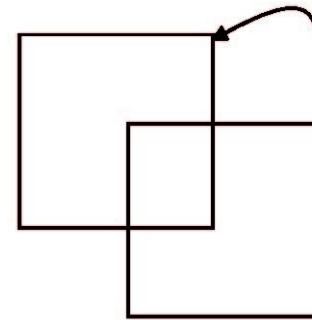
- OS sets instruction pointer to first instruction of process, non-memory-resident -> page fault

Issue: a given instruction could access multiple pages -> multiple page faults

- Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory
- Pain decreased because of **locality of reference**

Critical case: block move

- auto increment/decrement location
- Restart the whole operation?
- What if source and destination overlap?



Hardware support needed for demand paging

- Page table with valid / invalid bit
- Secondary memory (swap device with **swap space**)
- Instruction restart

Demand Paging

Worst case (access to page not in RAM)

- Context switch to OS (Page Fault trap, save state PRIOR instruction execution)
- Check that the page reference was legal and determine the location of the page on the disk
- Issue a read from the disk to a free frame:
 - Wait in a queue for this device until the read request is serviced
 - Wait for the device seek and/or latency time
 - Begin the transfer of the page to a free frame
- While waiting, allocate the CPU to some other user
- Receive an interrupt from the disk I/O subsystem -> Context switch to OS
- Correct the page table and other tables to show page is now in memory
- Switch back to faulting process

Performances:

- Measured with EAT
- Effective Access Time (EAT)
EAT =
 $(1 - p) \times \text{memory access}$
 $+ p (\text{page fault overhead} + \text{swap page in})$
 - Three major activities
 - Service the interrupt (~1k instructions), goes in overhead
 - Read/Write the page - lots of time
 - Restart the process - (~1k instructions), goes in overhead
- Page Fault Rate $0 \leq p \leq 1$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault

▪

Demand Paging Performance

Effective Access Time (EAT)

EAT =

$(1 - p)$ x memory access

+ p (page fault overhead

+ swap page in)

- Overhead (~hundreds of instructions):
 - Service the interrupt
 - Restart the process
 - Swap Page Out/In: lots of time
- p : page fault rate $0 \leq p \leq 1$
- if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault

Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$
$$= (1 - p) \times 200 + p \times 8,000,000$$
$$= 200 + p \times 7,999,800$$
- If one access out of 1,000 causes a page fault, then
$$EAT = 8.2 \text{ microseconds.}$$
Slowdown by a factor of 40!!
- If want performance degradation < 10 percent
 - $220 > 200 + 7,999,800 \times p$
 - $p < .0000025 \rightarrow (1/400,000)$

Demand Paging Optimizations

Swap space

Disk area without a file system (raw mode)

- I/O is faster than file system I/O even if it resides on the same device
 - NO filesystem overhead
- On startup: copy entire process image to swap space at process load time
- On execution: swap in and out of swap space
- When swapping out read only memory don't write back the data
- RW pages need to be written back when swapped out

Copy on Write (COW)

- When forking, replicate only the page table, to point to parent frames, but toggle a flag on the pages
- When forking, and set a "trap_on_write" flag on pages to 1
- on write a trap is generated
 - the frame is copied, and the bit is cleared so that further accesses will not trap
 - use reference counters on frames in OS to handle multiple forks

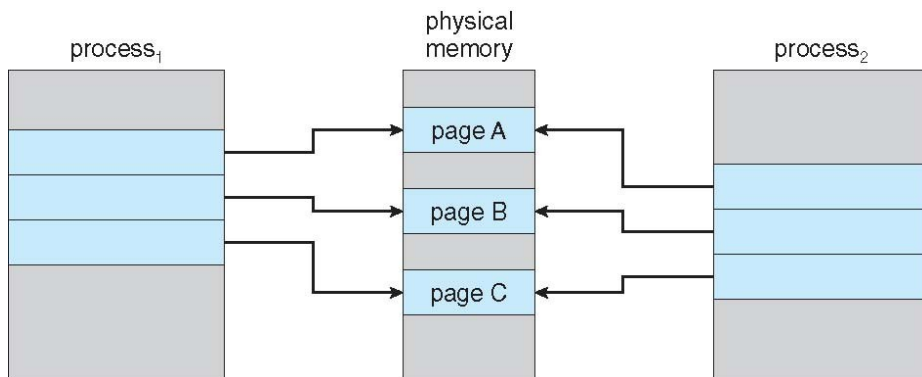
Free frames

- The system keeps a list of free frames (similar to a SLAB), to quickly get a free page when needed
- For security: free pages are zeroed (otherwise a new process might read the data of a dead process)

COW Example

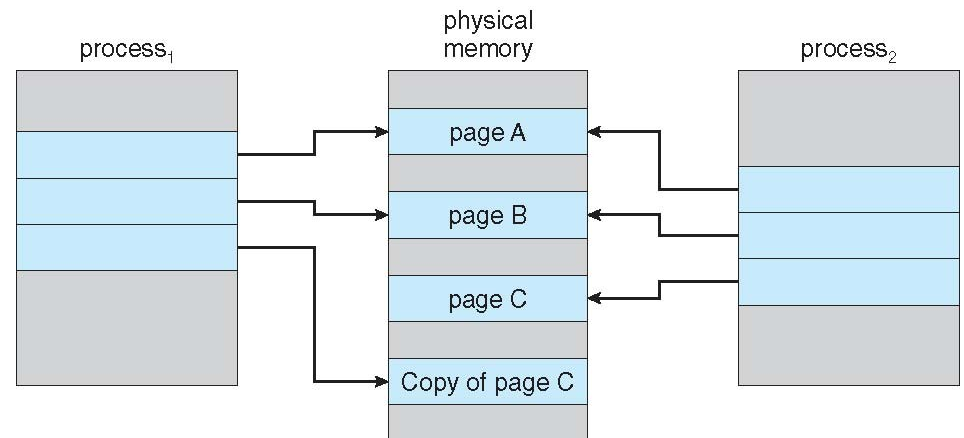
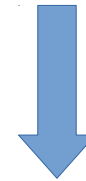
Before

- process 1 forks, and generates process 2
- page table is just copied, and trap_on_write bit is set
- frames are not copied



After

- when process 1 writes on page c, a trap is generated
- the frame "C" is copied, and the value in the page table of process 1 is updated



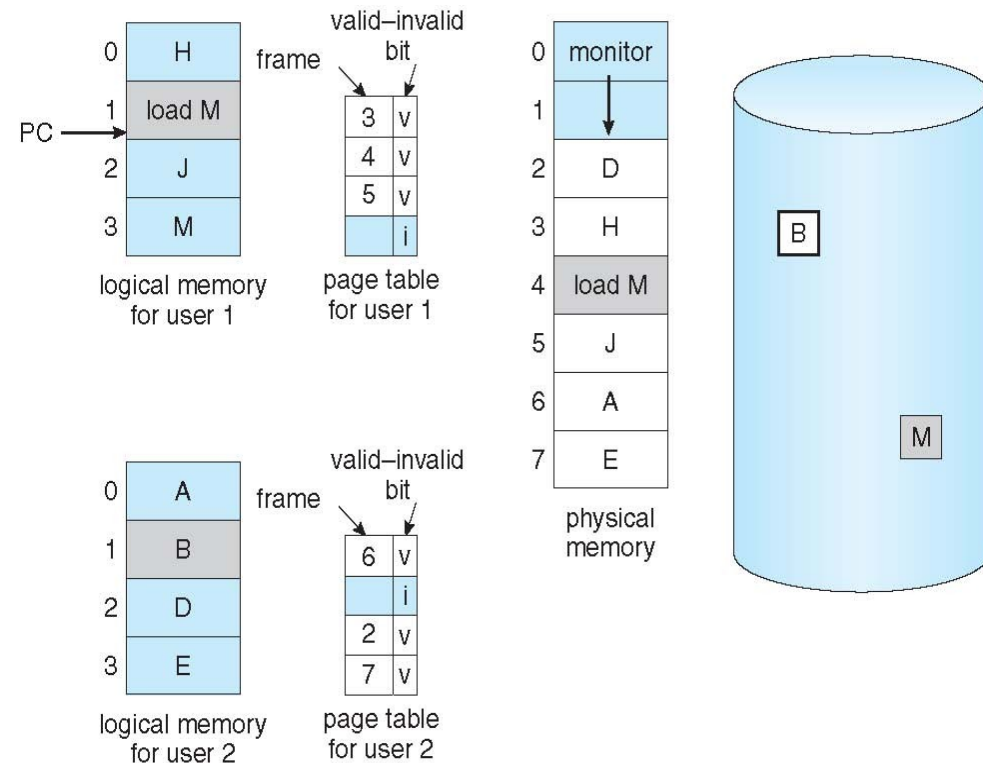
Page Replacement

Choose a used frame in memory to be swapped out (victim).

- Used when no free frame is available
- Optimality: choose the page that will be accessed latest
- requires knowledge about the future

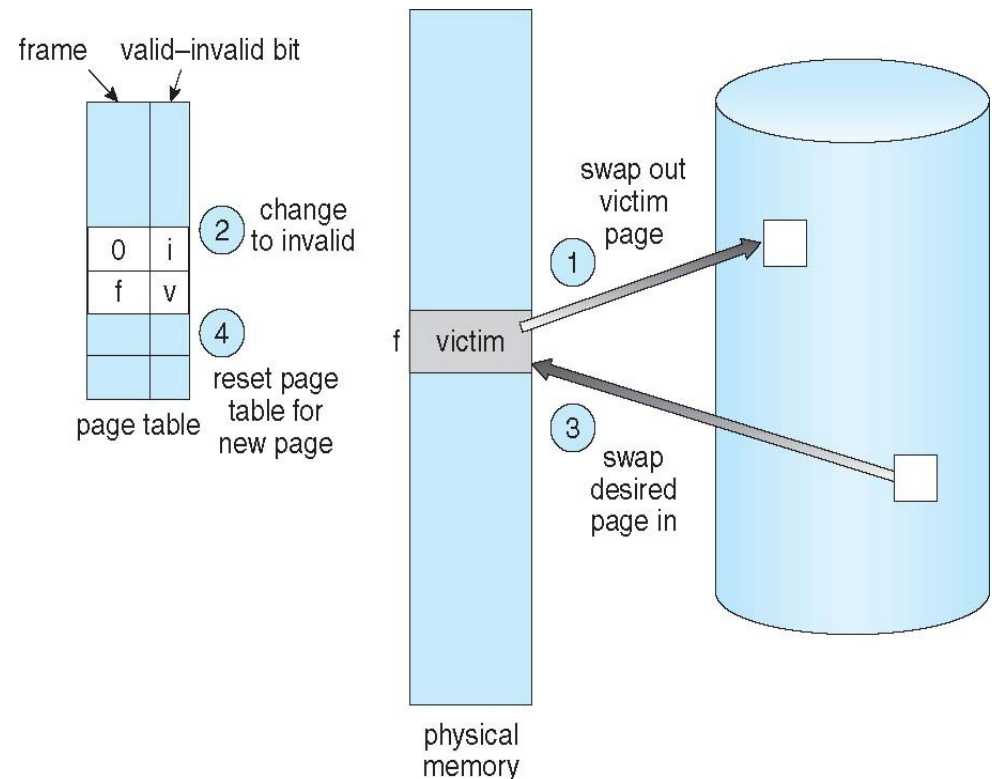
Pages that have not been altered in RAM do not need to be written back

- Use **modify (dirty) bit** to reduce overhead of page transfers
- They can be dismissed at lower cost



Basic Page Replacement

- Find the location of the desired page on disk
- Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if dirty
- Bring the desired page into the (newly) free frame; update the page and frame tables
- Continue the process by restarting the instruction that caused the trap



Note now potentially 2 page transfers for page fault – increasing EAT

Page Replacement Algorithms

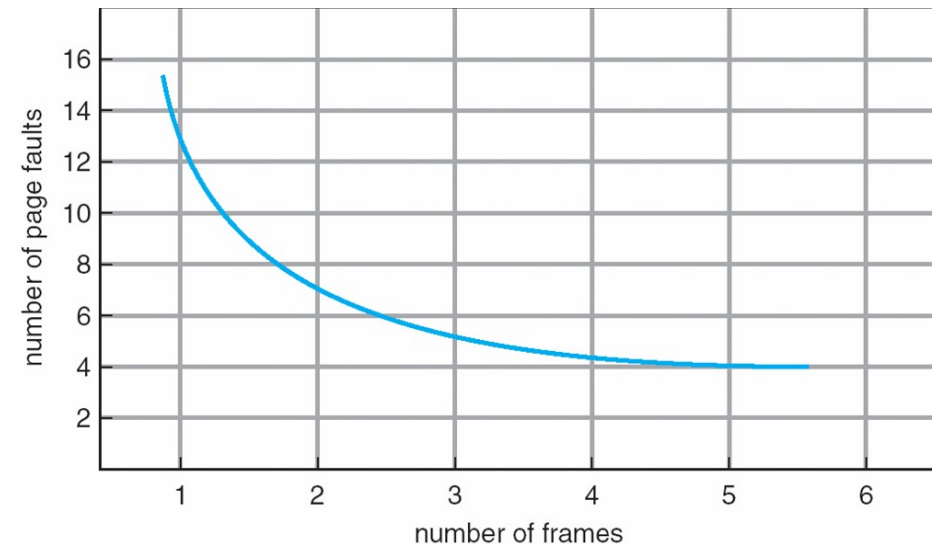
Goal: minimizing the page fault rate

- Evaluation:

- through simulation, using an array (string) encoding the access pattern
- $s[i]=x$, means that at time i , the system uses page x
- Example of reference string:

**<7,0,1,2,0,3,0,4,2,3,0,3,0,
3,2,1,2,0,1,7,0,1>**

The more the frames,
the less the page faults



PR: FIFO algorithm

Idea: Choose as victim the page that was swapped in last

- Example (3 frames)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

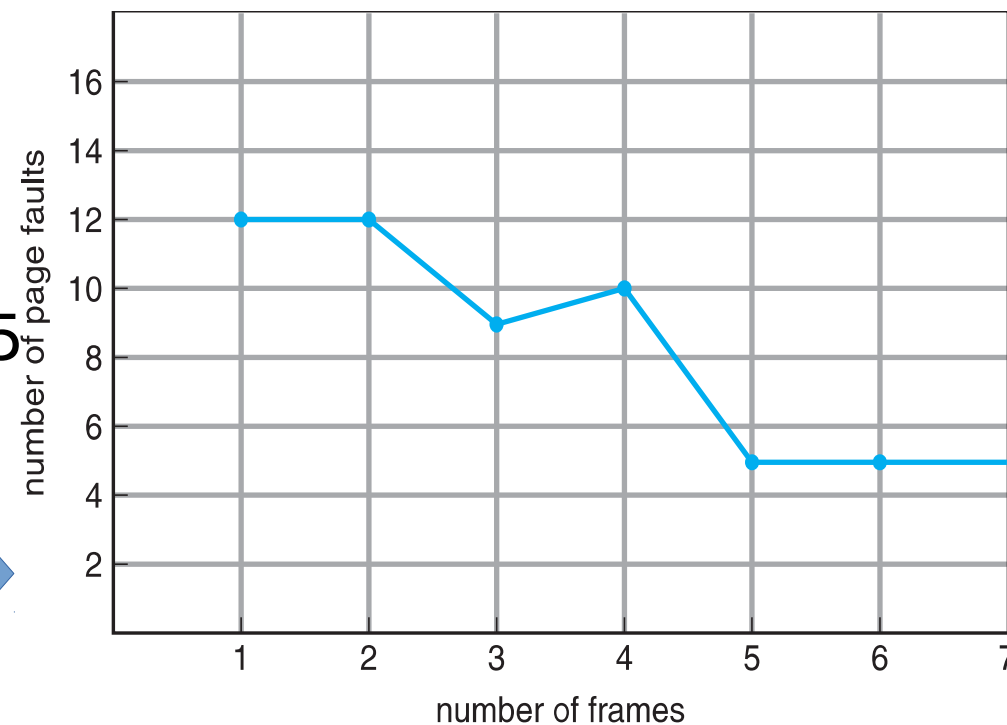
7	7	7	2														7	7	7
	0	0	0		2	2	4	4	4	0			0	0			1	0	0
		1	1		3	3	3	2	2	2			1	1			2	2	1

page frames

Can vary by reference string:
consider 1,2,3,4,1,2,5,1,2,3,4,5

- Adding more frames can cause more page faults!

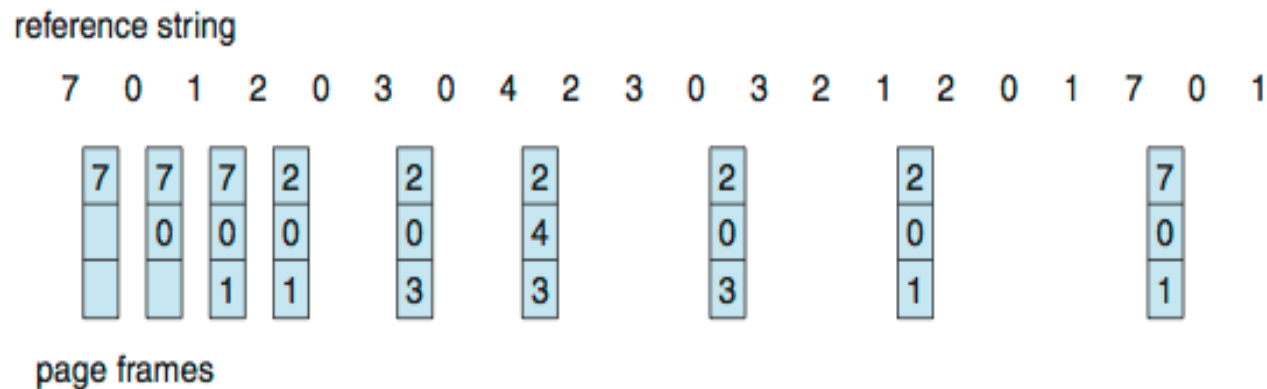
- Belady's Anomaly**



Optimal Page Replacement

Idea: Replace page that will not be used for longest period of time

- can't be done in practice
- provides an upper bound
 - all algorithms will be worse than optimal



Optimal value for this configuration is 9

LRU page replacement

Idea: approximate optimal by predicting which page will be used last.

Use prior knowledge to get the prediction: history repeats

Evicted page: the page that has not been used since longer

LRU not subject to Belady anomaly

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2				2		4	4	4	0				1		1		1
	0	0	0				0		0	0	3	3				3		0		0
		1	1				3		3	2	2	2				2		2		7

page frames

12 faults

LRU Implementations

- Counter:
 - each has a counter, when accessed copy the clock in the counter
 - on eviction: scan the page table
- List:
 - keep a list of pages. Each time a page is accessed, move its entry on top of the list.
 - expensive

Shortcomings: LRU requires special hardware, but it is still slow.

Full implementations not used.

Approximated implementations are.

LRU Approximations

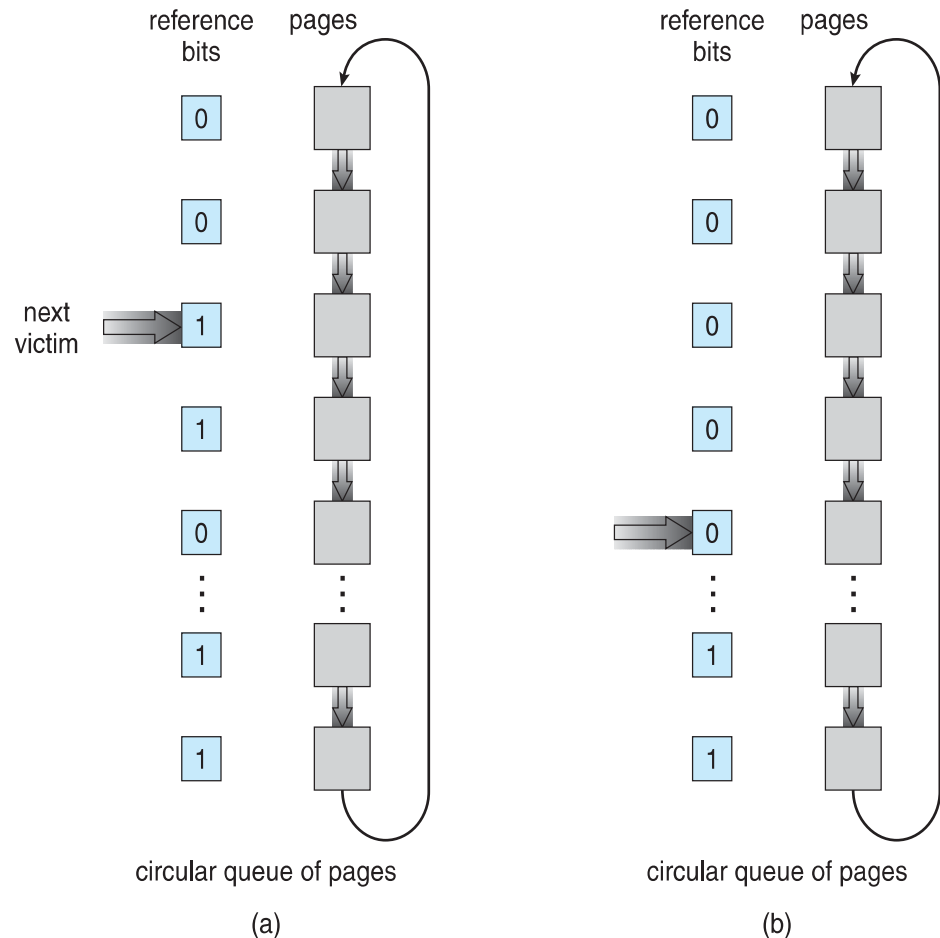
Need reference bit in page table
(HW support)

- With each page associate a bit, initially = 0
- When page is referenced bit set to 1
- Replace any with reference bit = 0 (if one exists)
 - We do not know the order, however

Second-chance algorithm

Clock replacement

- If page to be replaced has
 - Reference bit = 0 -> replace it
 - reference bit = 1 then:
 - set reference bit 0, leave page in memory
 - replace next page, subject to same rules



LRU approximations

Hardware Support: reference bit **and modify bit** in page table (HW support)

- When accessing a page set modify bit to 1

Enhanced Second-Chance Algorithm

Clock replacement

- rank pages based on access and modify bit
 - 0,0: best candidate (no write)
 - 0,1: write, but used long ago
 - 1,0: used recently, but no write
 - 1,1: worst case

Thrashing

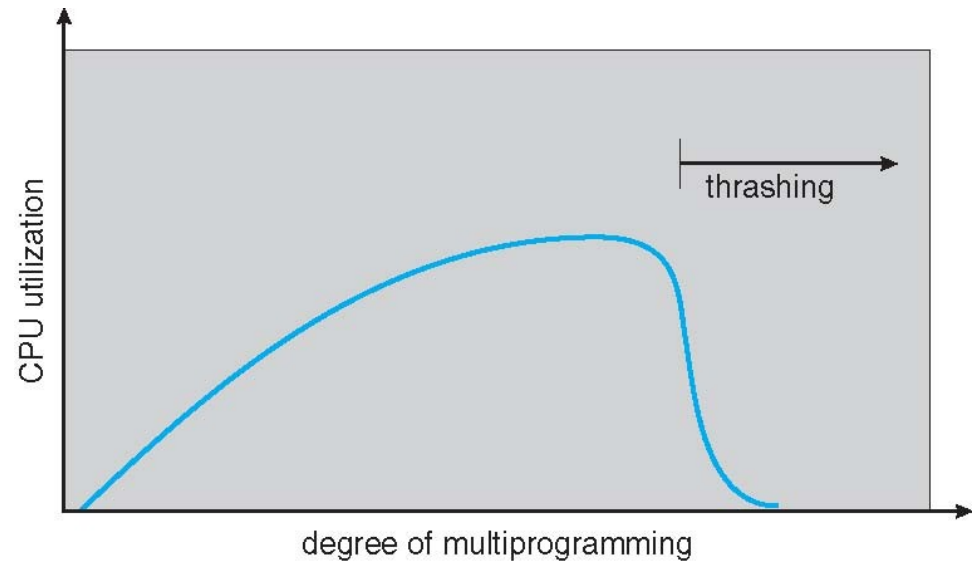
Thrashing: a process is busy swapping pages in and out

Happens when a process does not have “enough” pages, the page-fault rate is very high

- Page fault to get page
- Replace existing frame
- But quickly need replaced frame back

Consequences:

- Low CPU utilization
- Operating system thinking that it needs to increase the degree of multiprogramming
- Another process added to the system



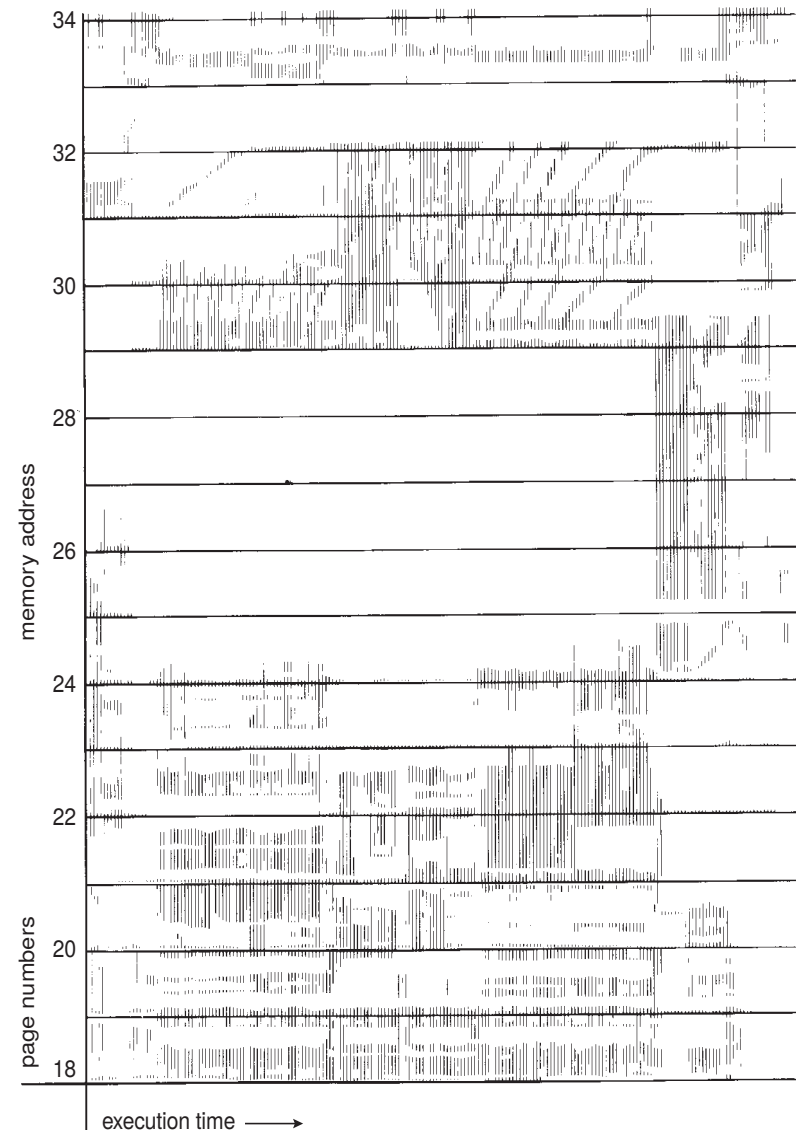
Memory Access Pattern

On the right we see the "pattern" of pages accessed as time evolves

Nearby columns show similar "black stripes": the regions of memory accessed as the system evolves changes smoothly

Locality principle:

- If I have accessed something short ago, it is very likely I will peek on it again in the near future



Working Set Model

Used to model access patterns and locality

Number of frames required at time t

$$D(t) = \sum_p w_{ss}(p, t)$$

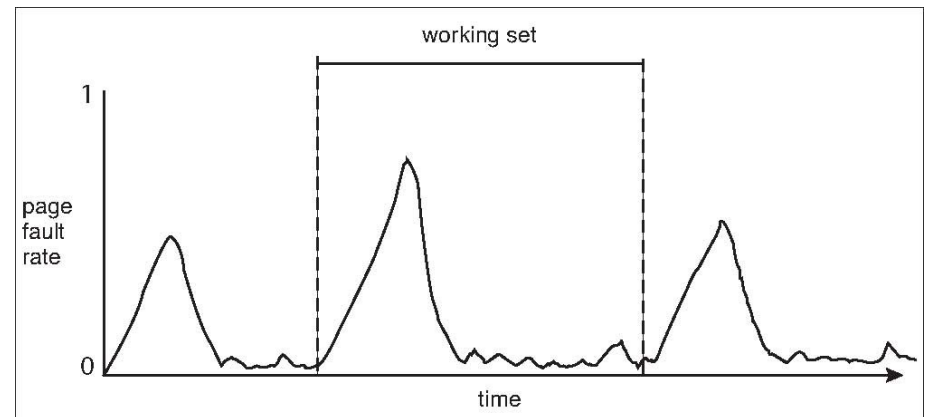
- $w_{ss}(p, t)$: function of two parameters:
 - t : the "epoch" (time interval) under analysis
 - p : the process id
 - $w_{ss}(p, t)$: set of pages accessed in the epoch t

Depends on the "duration" of an epoch

- if duration too small, not representative for the locality
- if duration too large, captures several localities

Trashing when too little frames available and too many required

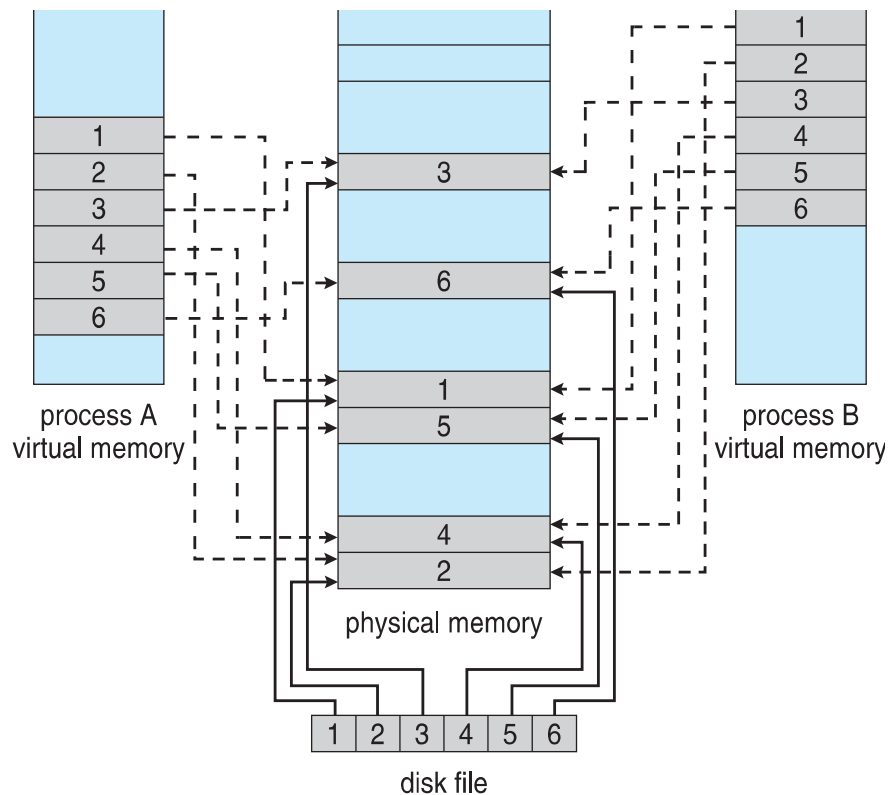
$$D(t) > \text{max_frames}$$



fault rate and working set are correlated

Constructs relying on VM

Memory mapped file
(mmap)



Shared Memory

