

Operating Systems

IPC

Giorgio Grisetti

`grisetti@diag.uniroma1.it`

Department of Computer Control and Management Engineering
Sapienza University of Rome

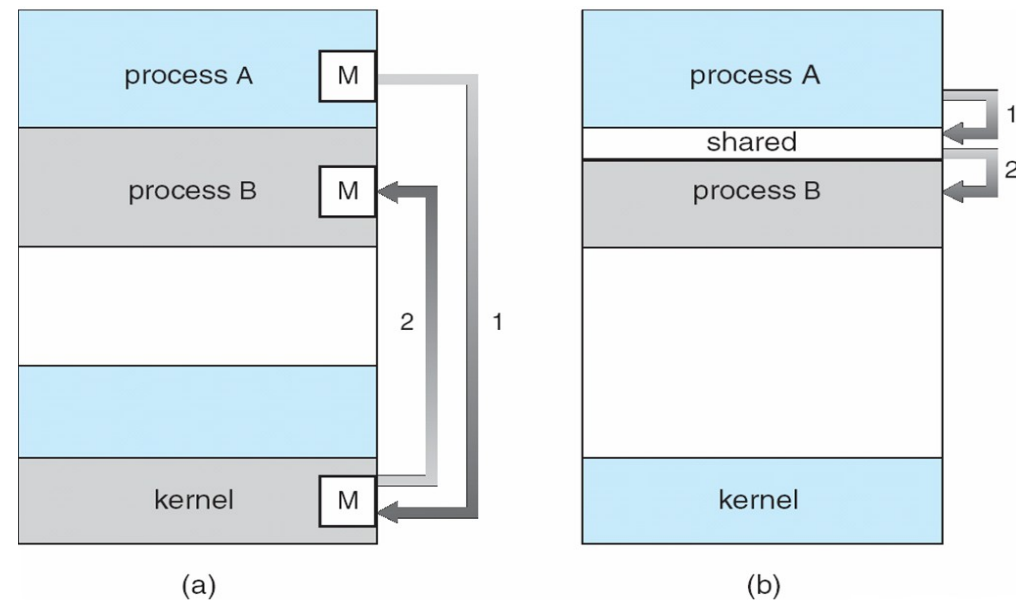
Interprocess Communication

In a complex system, processes might need to communicate with other processes

- sharing information
- enhancing the efficiency of the computation
- improving modularity
- cross check

Basic IPC

- message passing (a)
- shared memory (b)



Message Passing

Compared to shared memory, passing a message requires copying memory

- Pro: easier implementation
- Cons: decreased efficiency

Three operations

- `send(sender/mailbox, message)`
- `receive(sender/mailbox, message)`

Communication

- synchronous (blocking) vs asynchronous (non-blocking)
- direct (process) vs indirect (mailbox)
- limited vs unlimited

Message Passing

A message queue is a system-managed object that implements a mailbox

Processes need to know an identifier of the queue to operate on it

Functionalities

- open (creates or attaches to a message queue)
- close (detaches from a message queue)
- unlink (destroy)
- get/set attr (inspects the status of a queue/sets operation flags)
- put a message
- wait for a message/ notify when a message is ready

Posix message interface

- `mq_close()` close a message queue
- `mq_getattr()` get the current attributes of a message queue
- `mq_notify()` notify the calling process when the queue becomes nonempty
- `mq_open()` open or create a message queue
- `mq_receive()` receive a message from a queue
- `mq_send()` put a message into a message queue
- `mq_setattr()` set the flags for a message queue
- `mq_unlink()` unlink (i.e. delete) a message queue

Message Queue Example

```
#define QUEUE_NAME  "/test_queue"
#define MAX_SIZE    1024
#define MSG_STOP    "exit"
int run=1;
void * queueServer(void * args) {
    mqd_t mq;
    ssize_t bytes_read;
    struct mq_attr attr;
    char buffer[MAX_SIZE + 1];
    /* initialize the queue attributes */
    attr.mq_flags = 0;
    attr.mq_maxmsg = 10;
    attr.mq_msgsize = MAX_SIZE;
    attr.mq_curmsgs = 0;
    /* create the message queue */
    mq = mq_open(QUEUE_NAME,
                O_CREAT | O_RDONLY | O_NONBLOCK,
                0644, &attr);

    while(run) {
        memset(buffer, 0x00, sizeof(buffer));
        bytes_read = mq_receive(mq, buffer,
                                MAX_SIZE, NULL);

        if(bytes_read >= 0) {
            printf("SERVER: Received message: %s\n", buffer);
        } else {
            printf("SERVER: None \n");
        }
        fflush(stdout);
        usleep(0.725 * 1e6);
    }
    mq_close(mq);
    mq_unlink(QUEUE_NAME);
    return NULL;
}

void * queueClient(void * args) {
    mqd_t mq;
    char buffer[MAX_SIZE];
    mq = mq_open(QUEUE_NAME, O_WRONLY);
    int count = 0;
    while(run) {
        snprintf(buffer, sizeof(buffer),
                 "MESSAGE %d", count++);
        printf("CLIENT: Send message... \n");
        mq_send(mq, buffer, MAX_SIZE, 0);
        fflush(stdout);
        usleep(7.33 * 1e6);
    }
    mq_close(mq);
    return NULL;
}

int main(int argc, char** argv) {
    pthread_t client, server;
    printf("Start...\n");
    pthread_create(&server, 0, &queueServer, 0);
    pthread_create(&client, 0, &queueClient, 0);

    char c = getchar();
    run=0;
    pthread_join(server, NULL);
    pthread_join(client, NULL);

    printf("Done...\n");
    return (EXIT_SUCCESS);
}
```

Checking the queues

- From a shell, when the program is running you should see in the following directory
`/dev/mqueue` all queues created by the running processes and also those not unlinked
- A queue survives the creating process
- Doing `cat <queue_file>` displays the status of the queue

Shared Memory

- Is a memory area shared between processes
- Similar to the message queues or the semaphores, multiple shared memories can be present
- Functionalities are similar to those of processes
 - opening
 - closing
 - destroying
- A shared memory, once opened needs to be "mapped" to the memory area of a process through mmap

Shared Memory Writer

```
int main(int argc, char** argv) {
    // create a shared memory object at key argv[1]

    char* resource_name=argv[1];
    int fd=shm_open(resource_name, O_RDWR|O_CREAT, 0666);
    if (fd<0){
        printf("cannot create shared memory object, error: %s \n", sys_errlist[errno]);
        exit(-1);
    }
    int ftruncate_result = ftruncate(fd, SHMEM_SIZE);
    if (ftruncate_result<0) {
        printf("cannot truncate shared memory object, error: %s \n", sys_errlist[errno]);
        exit(-1);
    }
    // map the descriptor to the memory of the process
    void * my_memory_area = mmap(NULL, SHMEM_SIZE, PROT_WRITE, MAP_SHARED,
                                fd, 0);

    int num_rounds=100;
    for (int i=0; i<num_rounds; ++i) {
        char* buffer=(char*) my_memory_area;
        sprintf(buffer, "This is the message %d", i);
        printf("writing [%s]\n", buffer);
        sleep(1);
    }

    int unlink_result=shm_unlink(resource_name);
    if (unlink_result<0) {
        printf("cannot unlink shared memory object, error: %s \n", sys_errlist[errno]);
        exit(-1);
    }
}
```

Shared Memory Reader

```
int main(int argc, char** argv) {
    // create a shared memory object at key argv[1]

    char* resource_name=argv[1];
    int fd=shm_open(resource_name, O_RDONLY, 0666);
    if (fd<0){
        printf("cannot link to shared memory object, error: %s \n", sys_errlist[errno]);
        exit(-1);
    }
    int SHMEM_SIZE=0;

    struct stat shm_status;
    int fstat_result = fstat(fd, &shm_status); /* To obtain file size */
    if (fstat_result<0){
        printf("cannot get stats from memory object, error: %s \n", sys_errlist[errno]);
        exit(-1);
    }

    SHMEM_SIZE = shm_status.st_size;

    void * my_memory_area = mmap(NULL, SHMEM_SIZE, PROT_READ, MAP_SHARED,
                                fd, 0);

    while (1) {
        char* buffer=(char*) my_memory_area;
        printf("%s\n", buffer);
        usleep(100000);
    }
    int unlink_result=shm_unlink(resource_name);
    if (unlink_result<0) {
        printf("cannot unlink shared memory object, error: %s \n", sys_errlist[errno]);
        exit(-1);
    }
}
```

Checking the Shmem

- Similar to /dev/mqueue the shared memory is accessible through a virtual file in /dev/shm
- reading/writing to that file results in reading/writing to the shmem

Considerations

- We have seen the POSIX IPC API
- In conjunction to this API there is another set of functions that implement the SystemV primitives. These were prior the POSIX standardization
- In general the api is rather coherent for all objects
 - creation/opening: requires a string (filename) and returns an int (file descriptor)
 - specific functions operate on the object, solely based on the file descriptor
 - little things to remember, you can guess a lot

Exercises

- Implement a message queue by using a shared memory and semaphores